01 / The promise

One YAML. One verified agent.

Choose approved inputs. Let the operator configure the runtime and its connections. Prove the new agent before handing it to its user.

- The user selects model, MCP capabilities, skills, knowledge and an approved scenario.
- The operator and trusted platform owners create the per-agent runtime and Gateway security.
- Automatic authentication checks and the GoldenScenario produce a result tied to this exact request.

Office target, not a completed office result. Shared infrastructure, images and approved content are prerequisites.

02 / The architecture

The operator owns the wiring.

One owner request starts a controlled process. The agent runs as a normal approved workload without Substrate.

- Bind the actual ModelConfig, MCP tools and pinned skill/knowledge releases.
- Configure HTTPRoute, destination ReferenceGrants and strict JWT/authorization policies.
- Keep access fail-closed during reconciliation. Use independent bounded verification authority.
- Publish the usable endpoint only after the current security and GoldenScenario gates pass.

Architecture illustration. Real readiness comes from Kubernetes; actual calls come from observed telemetry.

03 / The live demonstration

Deploy, watch, invoke, verify.

Use a newly named agent so the audience can distinguish new onboarding from previously prepared resources.

- Apply the final owner YAML through the approved delivery process.
- Watch actual bindings, route/policy attachment and runtime readiness appear.
- Watch automatic JWT checks and GoldenScenario execution against the secured endpoint.
- Invoke as an approved caller, then show an unauthorized caller denied without forbidden dispatch.

No manual per-agent fixes or separately prepared scenario Job. Keep tokens, prompts and raw tool results private.

04 / The evidence view

See what is happening now.

Extend the existing read-only OTel view with onboarding and verification evidence. Do not build a replacement observability stack.

- Onboarding panel: named Kubernetes objects, current UID/generation, bindings and observation freshness.
- Execution panel: received model/Gateway/MCP spans, call names, status and duration; real skill/knowledge-use metadata.
- Verification panel: trusted JWT results, backend witnesses and the current GoldenScenario verdict.
- Unknown, blocked and stale stay visible. No timer-driven success or invented connections.

A received viewer trace does not prove Tempo ingestion. Verify traces, logs and metrics stores separately.

05 / What is already proved

Lab results and office gates stay separate.

The Gateway 1.7 home-lab receipts support a bounded technical result. The complete office onboarding demo still needs implementation and rehearsal.

- Lab: three in-scope Gateway containers upgraded to digest-pinned 1.7.0 images.
- Home lab, 3 November 2026: 30 security checks passed across two existing POCs, including allowed MCP calls and revocation checks.
- Lab: a real model request produced two received Gateway spans and one access log in the existing OTel viewer.
- Office still to prove: new one-file onboarding, automatic GoldenScenario, live lifecycle view and Tempo/log/metrics-store evidence.

Historical lab receipt included in the pack. It is not a full standalone-agent workflow, office deployment or Substrate result.

06 / The handover

Leave evidence the team can use.

At the end, present what was actually delivered and show how another engineer can repeat the same demonstration.

- Final owner YAML, editable-field guide and generated resource/binding inventory.
- Current GoldenScenario result with actual model, MCP, skill and knowledge evidence; PASS is required for completion.
- Protected live-view location, observed IDs, redacted captures and exact apply/invoke/watch/reset commands.
- Named operational owners, artifact prerequisites, rollback steps and explicit remaining limits.

Use the completion statement only after its gates pass. Substrate and Entra remain deferred; a Job must be described as a Job.